

# Singular Value Decomposition (SVD) Benchmarks

## 1. Introduction

### 1.1 SVD and Use Cases

Singular Value Decomposition (SVD) is a fundamental matrix factorization technique used widely in data science, machine learning, and signal processing. It decomposes a matrix  $A \in \mathbb{R}^{m \times n}$  into three matrices  $U$ ,  $\Sigma$ , and  $V^T$ , where:

- $U$  contains the left singular vectors,
- $\Sigma$  is a diagonal matrix with singular values,
- $V^T$  contains the right singular vectors.

SVD forms the backbone of various real-world applications including:

- Dimensionality reduction (e.g., Principal Component Analysis),
- Latent Semantic Analysis in natural language processing,
- Image compression and noise reduction in signal processing,
- Collaborative filtering in recommender systems.

Despite its versatility, SVD is computationally expensive, especially for large matrices. Thus, choosing an efficient implementation is crucial for scalable systems.

### 1.2 Goal of the Experiment

The objective of this experiment is to evaluate the performance of different SVD implementations in Python:

1. A naive raw Python implementation using basic operations,
2. NumPy's optimized `np.linalg.svd` function,
3. PyTorch's CPU-enabled `torch.linalg.svd` function.

The experiment aims to answer the following:

- How does each implementation perform in terms of execution time and scalability?
- How accurate is the reconstruction of the original matrix from its SVD components?
- What are the trade-offs between simplicity, speed, and accuracy?

## 2. Background

## 2.1 SVD Algorithm Brief Explanation

Given a matrix  $A \in \mathbb{R}^{m \times n}$ , its SVD is:  $A = U\Sigma V^T$  where:

- $U \in \mathbb{R}^{m \times m}$  is orthogonal (left singular vectors),
- $\Sigma \in \mathbb{R}^{m \times n}$  is diagonal (singular values),
- $V^T \in \mathbb{R}^{n \times n}$  is orthogonal (right singular vectors).

Computing SVD typically involves:

- Finding eigenvalues of  $A^T A$ ,
- Computing singular vectors via orthonormalization (e.g., Gram-Schmidt)
- Sorting components by descending singular values.

## 2.2 Differences in Implementations

- **Raw Python** manually performs the steps of SVD using basic Python operations and nested loops.
- **NumPy's `np.linalg.svd` function** leverages highly optimized linear algebra libraries such as LAPACK or OpenBLAS, which are written in C and Fortran. These libraries use advanced algorithms, efficient memory management, and can take advantage of multi-threading on modern CPUs.
- **PyTorch's `torch.linalg.svd` function** provides a similar high-performance SVD operation as NumPy. On CPU, PyTorch also relies on optimized linear algebra backends like MKL or OpenBLAS and supports multi-threaded execution.

## 3. Experiment Setup

### 3.1 Hardware and Software Environment

- CPU: Apple M2 Pro
- Total Number of Cores: 12
- GPU: None
- RAM: 16 GB
- OS: macOS
- Python 3.13.x
- NumPy 2.2.6
- PyTorch 2.7.x

### 3.2 Matrix Generation Method

To ensure reproducibility and control over experiment parameters, input matrices are generated using:

```
import numpy as np
```

```
def generate_random_matrix(n, m, seed=42):
```

```
np.random.seed(seed)
return np.random.rand(n, m)
```

- Square matrices ( $n=m$ ) such as 100x100, 200x200, 500x500, 1000x1000, 2000x2000 will be tested.
- The same matrix will be used across all implementations for fair comparison.

## 4. Methodology

### 4.1 Implementation

- Raw Python: A custom implementation using the power iteration method to compute the dominant (top-1) singular value and corresponding singular vectors.
- NumPy: Uses `np.linalg.svd`, internally calling optimized LAPACK routines.
- PyTorch: Uses `torch.linalg.svd`, suitable for large-scale data and deep learning pipelines.

Each implementation returns the top  $k$  singular values and vectors for fair benchmarking.

### 4.2 Metrics Definition

- Execution Time: Time required to compute the SVD of a single matrix (using `time.perf_counter`).
- Reconstruction Error: A measure of the accuracy of the decomposition, quantified by the Frobenius norm of the difference between the original matrix and its reconstruction from the computed SVD components.
- Scalability: An analysis of how execution time grows as the dimensions of the input matrix increase.
- Multi-processes Performance: An evaluation of how execution time is affected by varying the number of processes (1, 4, 8, 12) for a fixed matrix size.

All results will be reported with clear graphs and tables to highlight trade-offs between performance and accuracy.

## 5. Results Analysis

### 5.1 Single Process

#### 5.1.1 Performance Analysis

##### Expectation

- The optimized library implementations NumPy and PyTorch would significantly outperform the raw Python implementation in terms of execution time, especially

as the size of the input matrix increases. Both libraries use highly optimized low-level linear algebra routines (typically based on LAPACK/BLAS), while the raw Python implementation is not optimized and runs purely in Python, making it considerably slower.

- It was anticipated that NumPy and PyTorch would demonstrate similar performance on CPU for single-processed workloads, as both often rely on similar backend libraries.

## Result

```
In [1]: import pandas as pd
import numpy as np
import re
import matplotlib.pyplot as plt
import seaborn as sns
import glob

In [2]: # List all CSV files in the data folder
csv_files = glob.glob("data/benchmark_*.csv")

# Container for all DataFrames
dfs = []

# Process each file
for file in csv_files:
    # Extract parameters from filename using regex
    match = re.search(r"benchmark_s(\d+)_k(\d+)_i(\d+)_t(\d+)\.csv", file)
    if match:
        size, k, iterations, processes = match.groups()
        df = pd.read_csv(file)
        # Add metadata columns
        df['size'] = int(size)
        df['k'] = int(k)
        df['iterations'] = int(iterations)
        df['processes'] = int(processes)
        dfs.append(df)

# Concatenate all dataframes
merged_df = pd.concat(dfs, ignore_index=True)
print(merged_df.head())
```

	no	implementation	size	processes	time	frobenius_error \
0	1	raw_python	500	1	0.098452	275.258017
1	2	raw_python	500	1	0.098088	275.258017
2	3	raw_python	500	1	0.097933	275.258017
3	4	raw_python	500	1	0.097989	275.258017
4	5	raw_python	500	1	0.099348	275.258017

	relative_error	k	iterations
0	0.952962	1	100
1	0.952962	1	100
2	0.952962	1	100
3	0.952962	1	100
4	0.952962	1	100

```
In [3]: # Filter only data where processes == 1
single_process_df = merged_df[merged_df["processes"] == 1]
```

```

# Group by implementation and matrix size, then compute average time and
summary = single_process_df.groupby(["implementation", "size"]).agg({
    "time": "mean",
    "frobenius_error": "mean",
    "relative_error": "mean"
}).reset_index()

# Sort the results for easier interpretation
summary = summary.sort_values(by=["size", "implementation"])

implementation_order = ["raw_python", "numpy", "torch"]

summary["implementation"] = pd.Categorical(summary["implementation"], cat

summary = summary.sort_values(by=["size", "implementation"])

# Display the summary
print(summary)

```

	implementation	size	time	frobenius_error	relative_error
5	raw_python	100	0.003688	51.225233	0.895912
0	numpy	100	0.000894	28.525055	0.498894
10	torch	100	0.000884	28.525055	0.498894
6	raw_python	200	0.014467	106.769731	0.925931
1	numpy	200	0.003613	57.364574	0.497478
11	torch	200	0.003459	57.364574	0.497478
7	raw_python	500	0.097976	275.258017	0.952962
2	numpy	500	0.024499	144.055514	0.498730
12	torch	500	0.023895	144.055514	0.498730
8	raw_python	1000	0.410488	558.333066	0.966647
3	numpy	1000	0.117945	288.295932	0.499129
13	torch	1000	0.116556	288.295932	0.499129
9	raw_python	2000	1.696844	1127.424888	0.976379
4	numpy	2000	1.020877	577.027958	0.499721
14	torch	2000	1.011749	577.027958	0.499721

```

In [4]: # Define the exact sizes to plot
exact_sizes = [100, 200, 500, 1000, 2000]

# Filter the summary DataFrame to include only the exact sizes
filtered_summary = summary[summary["size"].isin(exact_sizes)]

# Create the plot
plt.figure(figsize=(10, 6))

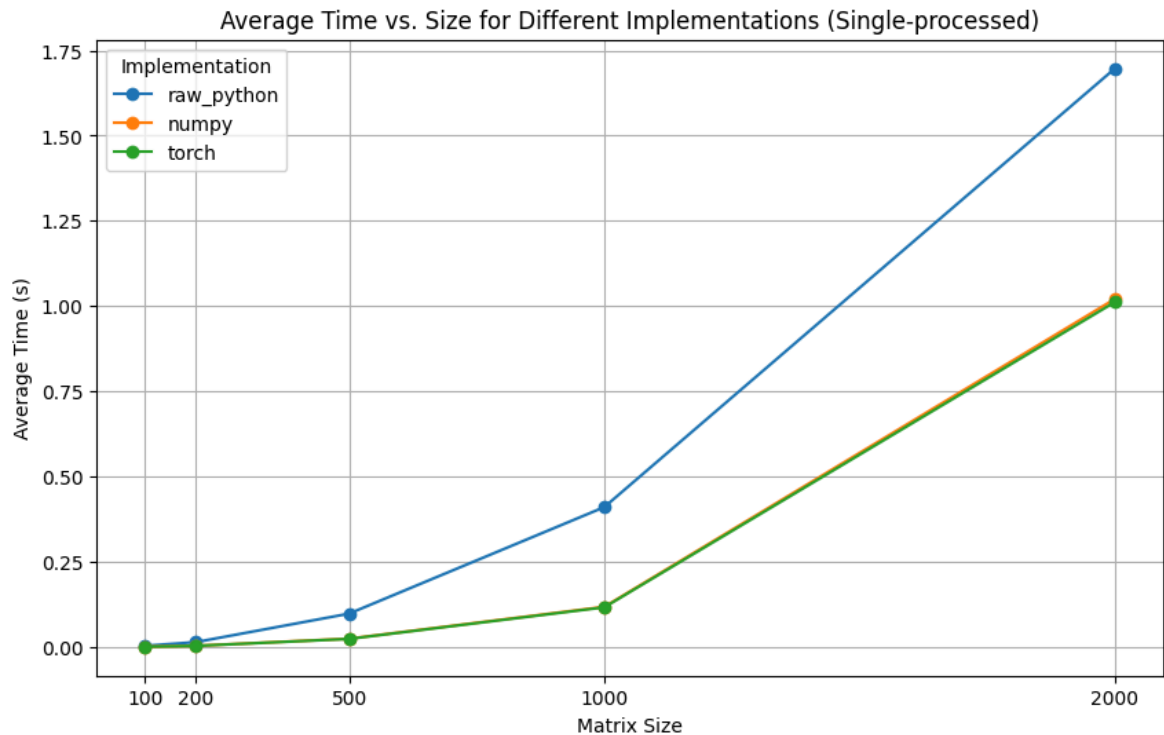
# Iterate through each implementation and plot the data
for implementation_name in filtered_summary["implementation"].unique():
    subset = filtered_summary[filtered_summary["implementation"] == implementation_name]
    plt.plot(subset["size"], subset["time"], marker='o', label=implementation_name)

# Add labels and title
plt.xlabel("Matrix Size")
plt.ylabel("Average Time (s)")
plt.title("Average Time vs. Size for Different Implementations (Single-process)")
plt.legend(title="Implementation")
plt.grid(True)

# Set x-axis ticks to show only the exact sizes
plt.xticks(exact_sizes)

```

```
plt.show()
```



## Conclusion

- Raw Python implementation is dramatically slower than both NumPy and PyTorch across all tested matrix sizes. As the matrix size increases, the execution time of raw Python grows much faster than that of the optimized libraries, highlighting its inefficiency for large-scale problems.
- NumPy and PyTorch perform almost identically in terms of execution time (PyTorch very slightly better), both being several times faster than raw Python. This suggests that for single-processed SVD on CPU, either library is an excellent choice for practical workloads.
- Overall, these results demonstrate the necessity of using optimized numerical libraries like NumPy or PyTorch for computationally intensive tasks such as SVD, particularly with larger datasets.

## 5.2.2 Accuracy Analysis

### Definition

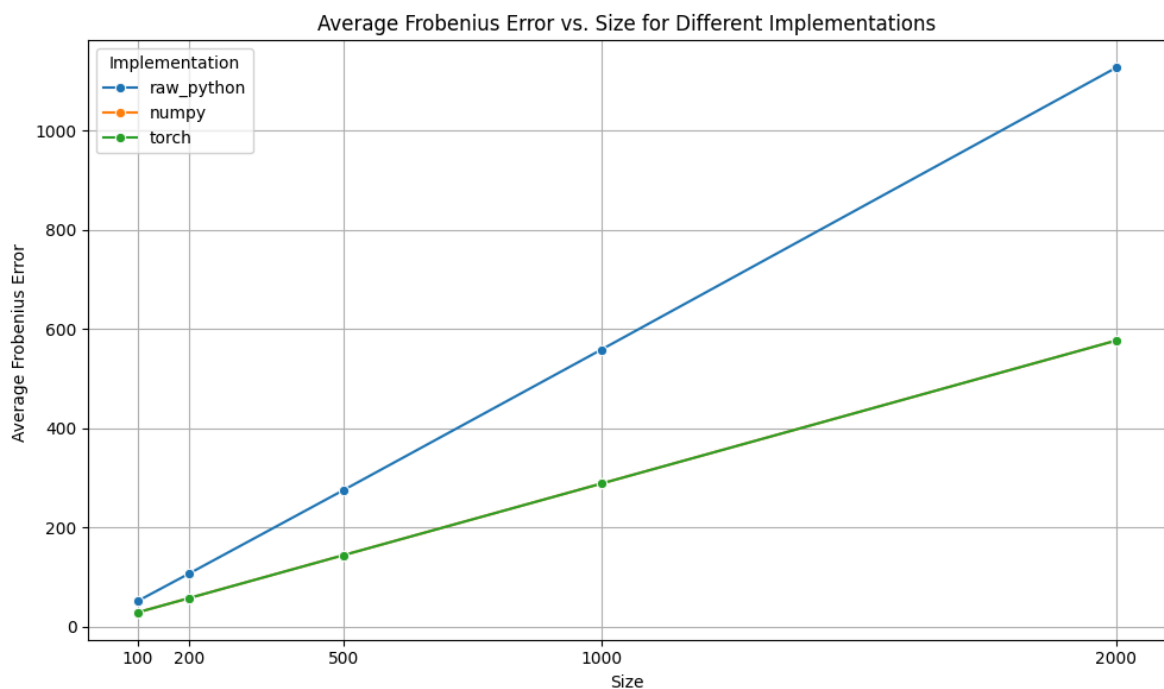
- Frobenius Error is a measure of how far the approximated matrix (from SVD) is from the original matrix. A lower Frobenius error means the approximation is more accurate.
- Relative Error is the Frobenius error scaled by the total "size" (norm) of the original matrix. It tells what percentage of the original matrix's information is lost. For example, relative error of 0.1 means 90% of the matrix is preserved.

### Expectation

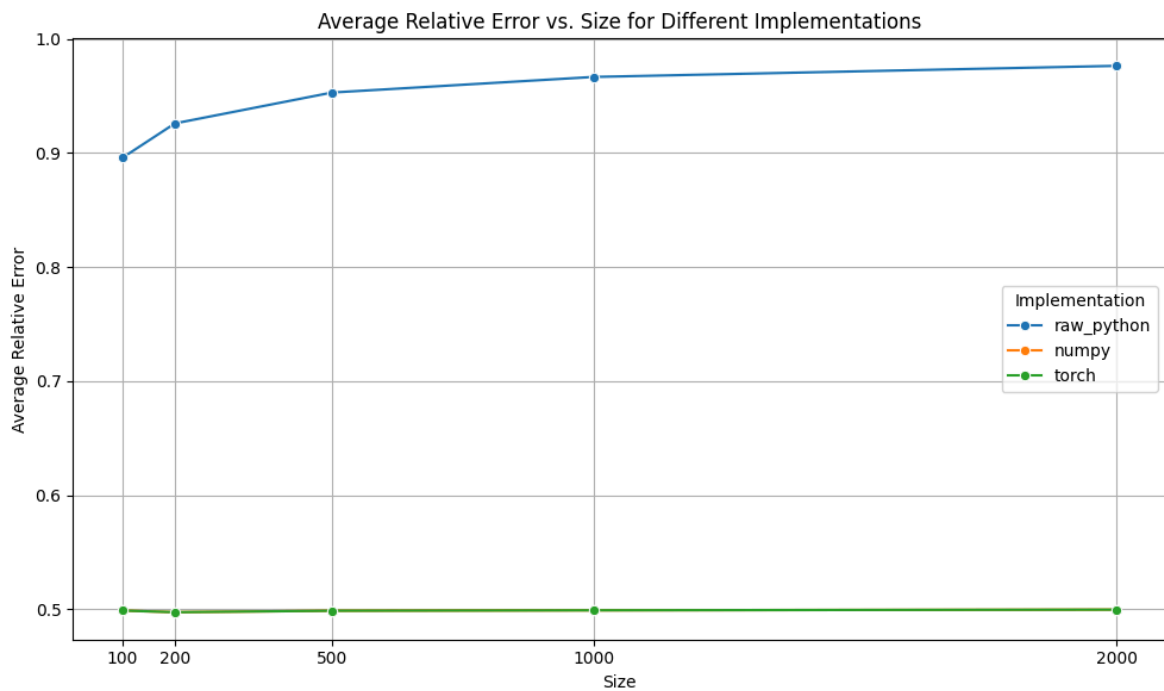
- NumPy and PyTorch implementations would achieve significantly lower Frobenius and relative reconstruction errors compared to the raw Python implementation. Since NumPy and PyTorch both use robust, well-tested linear algebra libraries for SVD computation, they should provide more accurate decompositions. The raw Python implementation, being a simple and less optimized algorithm (e.g., based on power iteration), was expected to show higher reconstruction errors, especially as the matrix size increases.
- It was anticipated that the accuracy of NumPy and PyTorch would be nearly identical due to their use of similar backend libraries.

## Result

```
In [5]: # Plotting the data
plt.figure(figsize=(10, 6))
sns.lineplot(data=summary, x='size', y='frobenius_error', hue='implementa
plt.title('Average Frobenius Error vs. Size for Different Implementations
plt.xlabel('Size')
plt.ylabel('Average Frobenius Error ')
plt.xticks(summary['size'].unique())
plt.grid(True)
plt.legend(title='Implementation')
plt.tight_layout()
plt.show()
```



```
In [6]: # Plotting the data
plt.figure(figsize=(10, 6))
sns.lineplot(data=summary, x='size', y='relative_error', hue='implementat
plt.title('Average Relative Error vs. Size for Different Implementations'
plt.xlabel('Size')
plt.ylabel('Average Relative Error')
plt.xticks(summary['size'].unique())
plt.grid(True)
plt.legend(title='Implementation')
plt.tight_layout()
plt.show()
```



## Conclusion

- NumPy and PyTorch both consistently achieve much lower Frobenius and relative errors than the raw Python implementation across all tested matrix sizes.
- For both error metrics, NumPy and PyTorch give the same result which indicates a high degree of consistency and accuracy between these two libraries. In addition, the average relative error for both libraries remain the same despite the matrix size differences.
- The raw Python implementation, by contrast, has much higher reconstruction errors, and these errors increase rapidly with the matrix size. The relative error for raw Python approaches 1.0 for large matrices, indicating that its reconstructions are far less accurate than those of NumPy or PyTorch.
- These results clearly demonstrate that for applications requiring accurate SVD decompositions, it is critical to use optimized libraries like NumPy or PyTorch, as naive implementations may not only be slower but also substantially less accurate.

## 5.2 Multi-Processing

### 5.2.1 Performance Analysis

#### Expectation

- Increasing the number of processes would lead to decreased execution times for SVD computations, especially for the optimized NumPy and PyTorch implementations, which are designed to leverage multi-processed linear algebra backends (such as OpenBLAS or MKL). The speedup was anticipated to be more noticeable for larger matrix sizes, as there is more work to parallelize.
- For the raw Python implementation, minimal or no performance improvement was expected with additional processes, since it is not parallelized and does not



utilize multi-processed libraries.

## Result

```
In [7]: sizes = [100, 200, 500, 1000, 2000]

# Define the process counts to analyze
multi_process_counts = [1, 4, 8, 12]

# Filter data for multiple processes
multi_process_df = merged_df[merged_df["processes"].isin(multi_process_co

# Group by implementation, matrix size, and processes, then compute avera
multi_process_summary = multi_process_df.groupby(["implementation", "size
    "time": "mean",
    "frobenius_error": "mean",
    "relative_error": "mean"
}).reset_index()

# Define the desired order for 'implementation'
implementation_order = ["raw_python", "numpy", "torch"]

# Convert 'implementation' to a categorical type with the desired order
multi_process_summary["implementation"] = pd.Categorical(multi_process_su

# Sort by the new categorical 'implementation' column along with size and
multi_process_summary = multi_process_summary.sort_values(by=["size", "pr

print("--- Performance Summary for Multiple processes (1, 4, 8, 12) ---")
print(multi_process_summary)

# Plotting Average Time vs. processes for each Implementation and Size
plt.figure(figsize=(15, 10))

# Get unique sizes for subplots
unique_sizes = multi_process_summary["size"].unique()
num_rows = (len(unique_sizes) + 1) // 2 # 2 plots per row

for i, s in enumerate(unique_sizes):
    plt.subplot(num_rows, 2, i + 1) # Create subplots
    size_subset = multi_process_summary[multi_process_summary["size"] ==

        for impl in implementation_order:
            impl_subset = size_subset[size_subset["implementation"] == impl]
            plt.plot(impl_subset["processes"], impl_subset["time"], marker='o

            plt.xlabel("Number of processes")
            plt.ylabel("Average Time (s)")
            plt.title(f"Size: {s}")
            plt.xticks(multi_process_counts)
            plt.legend(title="Implementation")
            plt.grid(True)
            plt.yscale('log')

plt.tight_layout()
plt.suptitle("Average Time vs. processes for Different Implementations an
plt.show()
```

--- Performance Summary for Multiple processes (1, 4, 8, 12) ---

	implementation	size	processes	time	frobenius_error	relative_e
rror						
20	raw_python	100	1	0.003688	51.225233	0.89
5912						
0	numpy	100	1	0.000894	28.525055	0.49
8894						
40	torch	100	1	0.000884	28.525055	0.49
8894						
21	raw_python	100	4	0.003868	51.225233	0.89
5912						
1	numpy	100	4	0.001081	28.525055	0.49
8894						
41	torch	100	4	0.001199	28.525055	0.49
8894						
22	raw_python	100	8	0.004087	51.225233	0.89
5912						
2	numpy	100	8	0.001159	28.525055	0.49
8894						
42	torch	100	8	0.001200	28.525055	0.49
8894						
23	raw_python	100	12	0.005347	51.225233	0.89
5912						
3	numpy	100	12	0.001528	28.525055	0.49
8894						
43	torch	100	12	0.001417	28.525055	0.49
8894						
24	raw_python	200	1	0.014467	106.769731	0.92
5931						
4	numpy	200	1	0.003613	57.364574	0.49
7478						
44	torch	200	1	0.003459	57.364574	0.49
7478						
25	raw_python	200	4	0.014932	106.769731	0.92
5931						
5	numpy	200	4	0.004444	57.364574	0.49
7478						
45	torch	200	4	0.004263	57.364574	0.49
7478						
26	raw_python	200	8	0.014962	106.769731	0.92
5931						
6	numpy	200	8	0.004773	57.364574	0.49
7478						
46	torch	200	8	0.004961	57.364574	0.49
7478						
27	raw_python	200	12	0.018191	106.769731	0.92
5931						
7	numpy	200	12	0.005345	57.364574	0.49
7478						
47	torch	200	12	0.005241	57.364574	0.49
7478						
28	raw_python	500	1	0.097976	275.258017	0.95
2962						
8	numpy	500	1	0.024499	144.055514	0.49
8730						
48	torch	500	1	0.023895	144.055514	0.49
8730						
29	raw_python	500	4	0.100379	275.258017	0.95
2962						
9	numpy	500	4	0.032564	144.055514	0.49

8730						
49	torch	500	4	0.030877	144.055514	0.49
8730						
30	raw_python	500	8	0.101574	275.258017	0.95
2962						
10	numpy	500	8	0.039637	144.055514	0.49
8730						
50	torch	500	8	0.037293	144.055514	0.49
8730						
31	raw_python	500	12	0.123207	275.258017	0.95
2962						
11	numpy	500	12	0.050897	144.055514	0.49
8730						
51	torch	500	12	0.046402	144.055514	0.49
8730						
32	raw_python	1000	1	0.410488	558.333066	0.96
6647						
12	numpy	1000	1	0.117945	288.295932	0.49
9129						
52	torch	1000	1	0.116556	288.295932	0.49
9129						
33	raw_python	1000	4	0.418993	558.333066	0.96
6647						
13	numpy	1000	4	0.163733	288.295932	0.49
9129						
53	torch	1000	4	0.163791	288.295932	0.49
9129						
34	raw_python	1000	8	0.421942	558.333066	0.96
6647						
14	numpy	1000	8	0.256933	288.295932	0.49
9129						
54	torch	1000	8	0.291517	288.295932	0.49
9129						
35	raw_python	1000	12	0.524564	558.333066	0.96
6647						
15	numpy	1000	12	0.432918	288.295932	0.49
9129						
55	torch	1000	12	0.422690	288.295932	0.49
9129						
36	raw_python	2000	1	1.696844	1127.424888	0.97
6379						
16	numpy	2000	1	1.020877	577.027958	0.49
9721						
56	torch	2000	1	1.011749	577.027958	0.49
9721						
37	raw_python	2000	4	1.732653	1127.424888	0.97
6379						
17	numpy	2000	4	2.595903	577.027958	0.49
9721						
57	torch	2000	4	2.576246	577.027958	0.49
9721						
38	raw_python	2000	8	1.729433	1127.424888	0.97
6379						
18	numpy	2000	8	7.004416	577.027958	0.49
9721						
58	torch	2000	8	7.142556	577.027958	0.49
9721						
39	raw_python	2000	12	2.151044	1127.424888	0.97
6379						
19	numpy	2000	12	10.118687	577.027958	0.49

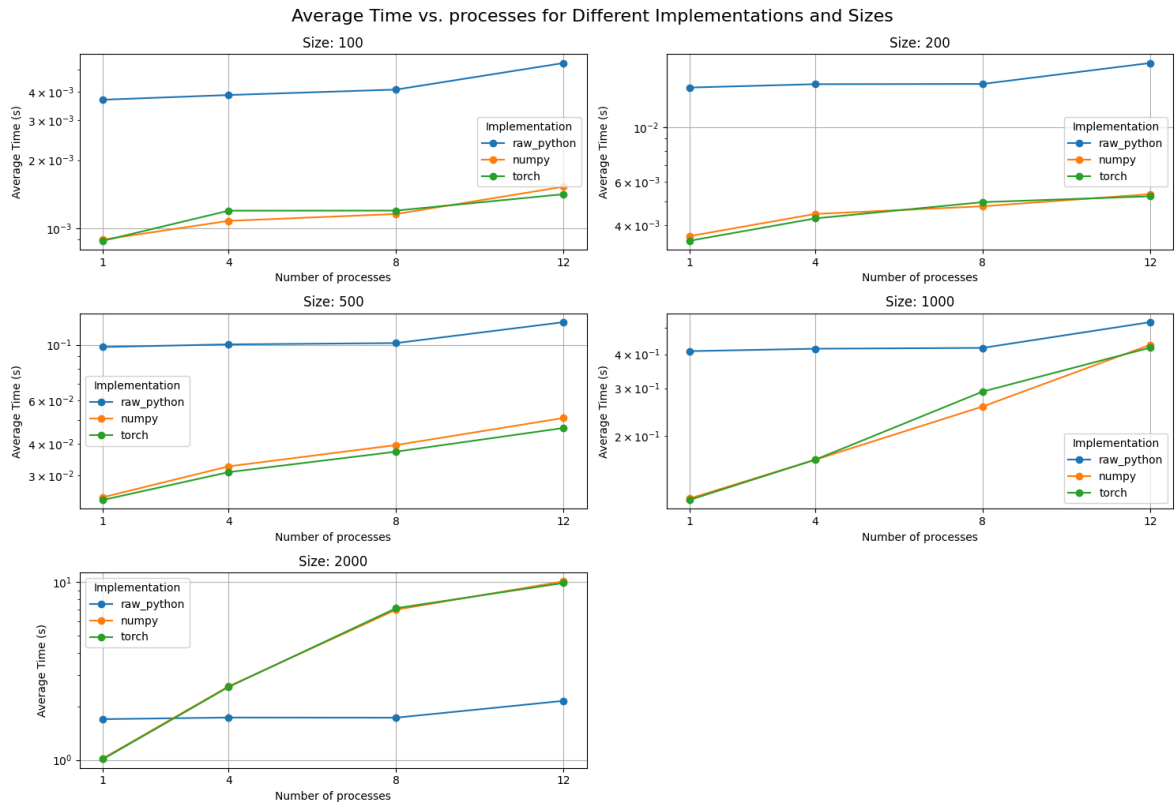
9721  
59  
9721

torch 2000

12 9.899629

577.027958

0.49



## Conclusion

- The data shows that for all three implementations (raw\_python, numpy, and torch) and across all tested matrix sizes (from 100x100 to 2000x2000), increasing the number of processes consistently led to an increase in average execution time. The performance did not improve with added processes; it degraded.
- This unexpected behavior may be caused by several factors:
  - **Raw Python**
    - No True Parallelism: Because of the Global Interpreter Lock (GIL), even on a multi-core processor, multiple processes in a CPU-bound Python program cannot run in parallel. They are forced to run sequentially.
    - Added Overhead: When increasing the number of processes, it is not gaining any parallel computation. Instead, it adds the overhead of context switching. The operating system has to spend time switching between the processes, each of which is just waiting for its turn to acquire the GIL.
  - **Numpy and Pytorch**
    - These libraries bypass the GIL by calling highly optimized and pre-compiled C, C++, or Fortran code (like OpenBLAS, MKL, or LAPACK). These underlying libraries can and do use multiple processes for computation.
    - However, the degrade in performance can be due to the cost of parallelization overhead which includes:

- **Process Creation & Management:** The process of creating, synchronizing, and tearing down processes takes time.
- **Data Partitioning & Aggregation:** The matrix must be broken down into smaller pieces for each process to work on, and the results must be combined afterward. This is not an instantaneous process.
- **Memory Bandwidth & Cache Issues:** SVD is a memory-intensive operation. When multiple processes run simultaneously, they compete for access to the system's memory bus and CPU caches. This contention can lead to processes waiting for data, negating the benefits of parallel processing.
- For the matrix sizes in this experiment, the computation itself is likely so fast on modern CPUs that the overhead of managing the processes is greater than the time saved by parallel execution. It is simply faster for a single core to perform the entire calculation than it is to coordinate the work among multiple cores. This is common for tasks that are not sufficiently large or complex to amortize the cost of parallelization.
- The experiment demonstrates a critical concept in performance engineering: more processes does not automatically mean better performance. It highlights that blindly adding processes without considering the problem size and the overhead involved can be counterproductive.

## 6. Limitations

- **Hardware Constraints:** All experiments were conducted on a single hardware configuration (Apple M2 Pro, CPU only). Results may differ on systems with different CPUs, more cores, or on dedicated GPUs, especially for PyTorch.
- **Fixed Data Size Range:** Matrix sizes tested were between 100 and 2000. Larger or smaller matrices, or non-square matrices, may result in different trends.

## 7. Conclusion

This study systematically benchmarked three different SVD implementations: raw Python, NumPy, and PyTorch, under varying matrix sizes and process counts to evaluate their performance and accuracy. The results confirm that the choice of implementation is critical. Both NumPy and PyTorch consistently delivered dramatically faster and more accurate SVD computations compared to the naive `raw_python` implementation. This highlights the indispensable role of optimized, low-level backend routines (like BLAS and LAPACK) in numerical computing.

For any real-world application requiring SVD, using optimized libraries like NumPy or PyTorch is non-negotiable for achieving acceptable performance and accuracy. However, this study serves as a critical reminder that "more processes" does not automatically equate to "more speed." Performance benefits from multi-processing are only realized when the computational task is large enough to amortize the

associated overhead. Therefore, developers should be cautious about parallelization strategies and recognize that for moderately-sized problems, a single-processed approach may yield the best performance.